Replayable Agent Evaluation Sandbox

Company inspiration: Halluminate | End-to-end build instructions + copy/paste coding-agent prompts

01 | Product brief

Problem: Long-running knowledge-work agents are hard to train, evaluate and reproduce.

Build: Disposable browser/spreadsheet task worlds with deterministic grading, snapshots and replayable action traces.

Audience: engineering teams and technical evaluators. Deliverable: working local MVP, reproducible tests, benchmark report, short demo and documented limitations.

This is an independent proof-of-work proposal. It does not assert knowledge of the company's private architecture or current incidents.

02 | Recommended architecture and repository

Java 21 + Spring Boot 3 for APIs/orchestration, React + TypeScript for dashboard, PostgreSQL for state, Redis for short-lived state, Kafka where asynchronous events are justified, Docker Compose locally, AWS/Kubernetes later. Substitute Python services where ML/browser libraries require it.

Monorepo: /apps/web, /services/api, /services/worker, /packages/contracts, /infra/docker, /tests/e2e, /benchmarks, /docs. Use OpenAPI, structured logs, OpenTelemetry, health checks and environment-variable configuration.

03 | Build milestones (execute in order)

Phase 0 - Contract and safety — Write README, requirements, threat model, architecture diagram, non-goals and measurable acceptance criteria. Create repository, .gitignore, CI and coding conventions.

Phase 1 - Working skeleton — Bootstrap backend, frontend, PostgreSQL and local Docker Compose. Add health endpoints, migrations, request validation and one end-to-end smoke test.

Phase 2 - Core data and APIs — Implement domain entities, API contracts, permissions, idempotency, state transitions and integration tests.

Phase 3 - Functional engine — Implement the domain-specific modules below, one by one, with unit tests and failure handling.

Phase 4 - Agent and automation — Introduce optional coding/AI agents only behind explicit interfaces, limits, trace logging and deterministic fallbacks.

Phase 5 - Reliability and observability — Add retries, dead-letter handling where needed, timeouts, rate limits, dashboards, tracing and chaos/failure scenarios.

Phase 6 - Evaluation and packaging — Run benchmarks, write limitations, record demo, prepare deployment guide and clean sample dataset.

04 | Functional modules - implementation checklist

- Define task DSL, fixtures and expected end-state assertions
- Create sandbox world with seeded browser and spreadsheet app
- Agent runner with explicit tool contracts and budgets
- Capture action/event trace and world snapshots
- Deterministic graders, replay UI and curriculum generator

05 | API, persistence and events

Representative API routes: POST /tasks, POST /runs, GET /runs/{id}/trace, POST /runs/{id}/replay

Core tables: task_templates, worlds, runs, actions, snapshots, grades

Async events (only if needed): run.started, action.recorded, snapshot.created, run.graded

Design rules: every mutation validates input and permissions; writes are idempotent where retried; DB migrations are reversible where practical; event consumers are at-least-once safe; every state transition is observable.

06 | Copy/paste agent prompt: start from zero

```
You are the lead staff engineer and implementation agent for an ORIGINAL prototype called "Replayable Agent Evaluation Sandbox" inspired by the public product area of Halluminate. Do not imply affiliation or access to private company systems.
MISSION: Long-running knowledge-work agents are hard to train, evaluate and reproduce.
BUILD: Disposable browser/spreadsheet task worlds with deterministic grading, snapshots and replayable action traces.
STACK: Java 21 + Spring Boot 3 for APIs/orchestration, React + TypeScript for dashboard, PostgreSQL for state, Redis for short-lived state, Kafka where asynchronous events are justified, Docker Compose locally, AWS/Kubernetes later. Substitute Python services where ML/browser libraries require it.
FIRST: inspect the repository. If empty, scaffold the monorepo. Write docs/PRD.md, docs/ARCHITECTURE.md, docs/THREAT_MODEL.md, docs/PLAN.md, and docs/DECISIONS.md. Specify user stories, non-goals, interfaces, threat model, and measurable acceptance tests.
THEN: implement in small vertical slices. After every slice, run formatting, compilation, unit tests, and smoke tests. Fix failures before proceeding. Commit logically and update docs/PLAN.md with completed work.
MANDATORY MODULES: Define task DSL, fixtures and expected end-state assertions; Create sandbox world with seeded browser and spreadsheet app; Agent runner with explicit tool contracts and budgets; Capture action/event trace and world snapshots; Deterministic graders, replay UI and curriculum generator.
CONTRACTS: POST /tasks, POST /runs, GET /runs/{id}/trace, POST /runs/{id}/replay. TABLES: task_templates, worlds, runs, actions, snapshots, grades. EVENTS: run.started, action.recorded, snapshot.created, run.graded.
CONSTRAINTS: Use fictional data and isolated tools; never connect agent evaluations to real accounts or money movement.
Never invent benchmark results or claim production readiness. Ask for approval before destructive actions, cloud spending, deploying, or accessing real user data. Provide exact commands to run locally. Start by writing the plan and scaffolding Phase 0 and Phase 1.
```

07 | Copy/paste agent prompt: implement core engine

```
Continue the Replayable Agent Evaluation Sandbox repository from the current state. Read docs/PLAN.md and tests before changing code. Implement the NEXT unfinished domain module, prioritizing: Define task DSL, fixtures and expected end-state assertions; Create sandbox world with seeded browser and spreadsheet app; Agent runner with explicit tool contracts and budgets; Capture action/event trace and world snapshots; Deterministic graders, replay UI and curriculum generator. Add schemas, migrations, API validation, permissions, retry semantics, structured logging, unit tests and at least one end-to-end happy-path and failure-path test. Keep changes reviewable; do not skip tests or silently change architecture. At the end report changed files, exact run commands, test results, unresolved risks and next step.
```

08 | Copy/paste agent prompt: testing, evaluation and hardening

```
Act as an independent QA, security and performance engineer for Replayable Agent Evaluation Sandbox. Do not assume the implementation works. Read the PRD, inspect code and execute tests. Create reproducible fixtures and benchmarks for: Same seeded task yields same expected grader result; Replay identifies the first divergence; Report task success, action count, cost and grader agreement. Test timeouts, retries, duplicate events, partial failures, permissions, injection risks, data leakage and worker crashes as applicable. Record real observed results, not fabricated numbers. Create docs/BENCHMARKS.md and docs/SECURITY_REVIEW.md with reproducible commands and prioritized fixes. Implement the highest-severity fixes and rerun tests. Enforce: Use fictional data and isolated tools; never connect agent evaluations to real accounts or money movement.
```

09 | Copy/paste agent prompt: demo and deployment

```
Prepare Replayable Agent Evaluation Sandbox as a public engineering proof-of-work. Add sample synthetic data, one-command local setup, seed script, health checks, demo walkthrough, architecture diagram, API documentation and troubleshooting. Add a 3-minute demo script that shows the problem, core workflow, failure recovery and measured results. Provide optional AWS deployment architecture and rough cost assumptions, but do not provision cloud resources without explicit approval. Ensure README distinguishes implemented features from future ideas.
```

10 | Definition of done and benchmark plan

• Same seeded task yields same expected grader result

• Replay identifies the first divergence

• Report task success, action count, cost and grader agreement

• Fresh clone can start locally using documented commands.

• Unit, integration and end-to-end tests pass in CI.

• Security restrictions and human approval boundaries are documented and tested.

• Benchmark numbers are reproduced from scripts and committed sample fixtures.

• README includes limitations, architecture trade-offs and a short product demo.

11 | Suggested four-sprint execution

Sprint 1: PRD, scaffolding, domain model, API skeleton and local stack.

Sprint 2: core modules, happy-path integration, synthetic data and dashboard.

Sprint 3: failure handling, security boundaries, telemetry and agent adapters.

Sprint 4: benchmarking, documentation, polish and demo recording.

12 | Next-level research directions

After the MVP works, identify the highest-cost or lowest-reliability component from measured data. Compare a simple baseline against one optimized algorithm. Publish an experiment with dataset, configuration, confidence intervals where meaningful, and limitations. Avoid unvalidated performance claims.